

INSIGHTS REPORT

Learning Rate: $1e-4$ (stable convergence, less oscillation)

Optimizer: Adam (faster convergence and better generalization)

Batch Size: 8 (balanced between stability and GPU memory)

Epochs: 50

Loss Function: MSELoss (worked best for pixel-level regression)

Augmentations: Random horizontal flip

MSELoss performed better than L1 as the task is to match RGB pixel values precisely. Adam outperformed SGD with momentum by adapting learning rates per parameter. Batch size 8 gave the best training speed vs memory usage trade-off.

UNet Architecture and Conditioning

Final UNet Design:

- **Input Channels:** 6
→ 3 for base image + 3 for polygon mask
- **Output Channels:** 3 (for RGB)
- **Depth:** 4 levels (Encoder → Bottleneck → Decoder)
- **Skip Connections:** Yes (standard in UNet)
- **Activation:** ReLU in intermediate layers, Sigmoid/Linear at output (to maintain RGB range)
- **Normalization:** BatchNorm

Ablations Tried:

- Tried input as single channel mask vs 3-channel → 3-channel preserved structure better.
- Tried deeper UNet (5 levels) → increased training time with minimal gains.
- Skip connection removal → led to blurrier results (spatial info loss).

Training Dynamics

Loss & Metrics:

- Training Loss dropped smoothly with $LR=1e-4$.
- Validation loss plateaued ~epoch 25.
- No signs of overfitting observed until ~30 epochs.

Qualitative Trends:

- Early epochs: Model learned outlines, colors were random.
- Mid training: Started matching general color regions.
- Final epochs: Captured object-specific colors and boundaries.

Failure Modes:

- Some shapes near image edges were under-colored.
- Rare class polygons (low frequency) were poorly predicted.
- Polygon overlaps caused color bleed occasionally.

Fixes Attempted:

- Added more balanced samples (underrepresented polygons).
- Applied stronger augmentations (flip, scaling).
- Tweaked polygon mask representation from binary to RGB channels for richness.

Key Learnings

1. Mask Conditioning Matters

Representing polygon masks as RGB helped preserve structural information better than binary masks.

2. Model Size Balance

A moderate-depth UNet worked well. Larger models didn't bring much gain, proving that smart input conditioning is more powerful than raw depth.

3. Data Quality > Model Complexity

Improving dataset quality (augmentations, mask detail) had more effect than changing architectures.

4. W&B Artifacts FTW

Using Weights & Biases for model versioning made collaboration and deployment easier, especially for sharing latest models directly.

5. Inference Notebook Should Be Lightweight

Separating training from inference helped keep things modular and clean.